

Jev *vs.* the Frontier

Same accuracy on
automated decisions.
One-fiftieth of the cost.

91.4% *vs* 91.0%

Accuracy on auto-accepted messages, Jev vs
Claude Sonnet 5


52×

lower cost per decision
(\$0.10 vs \$5.25 per 1,000)

Swipe for the findings →

Same accuracy. 52× cheaper.

Three models, one automated pipeline,
1,000 held-out bank messages. Only
confident answers get automated. How
accurate were they, and what did they
cost?



Accuracy on auto-accepted messages, 70% setting · cost per 1,000 (log scale)

Cheap first beat expensive alone.

Let Jev handle what it's sure about.
Send only the rest to Sonnet 5. **Every
message still gets an answer.**

73%

Jev auto-accepts · 91.4% right

27%

to Sonnet 5

Cascade: Jev first,
Sonnet for the rest

82.8%

accuracy · **\$1.51** / 1,000

Sonnet 5 on
every message


82.3%

accuracy · **\$5.25** / 1,000

Slightly more accurate for 29% of the cost. The best system wasn't one model. It was a routing design.

Jev's confidence can be dialed in.

Asked how sure they are, the Claude models mostly answer **85, 90 or 95**. So many answers tie that you **can't set the automation level precisely**.



Each dot is one automation level a team could choose (test set). Ring = what the "70%" setting actually automated.

Leaderboards are the wrong scoreboard.

- 01 **Pick models by cost per trusted decision.** Sonnet was slightly better on *all* answers, and would have looked like the winner.
- 02 **Confidence is a product feature.** It sets your automation rate and your error rate.
- 03 **System design moved the needle** as much as model choice: rule-based checks, a confidence bar and a cascade.

Full results, interactive charts and the technical report

danielwipert.github.io/jev.sorter

Pre-registered · 1,000 held-out messages · every log public · total cost \$16

